

Artificial Intelligence

AI2002

Rushda Muneer

Spring 2026

Course Information

- Instructor: Rushda Muneer
- Email: rushda.muneer@lhr.nu.edu.pk
- Office: F-024 (CS Department)
- Office Hours: Tuesday & Thursday (1:00 -2:00 pm)
- Google Classroom Codes:

BCS 6B (1-2:30)

xrxxzajo

BCS 6C (10-11:30)

etdpf7gs

General Rules and Guidelines

- Attendance will be marked after **10 minutes** in each class.
- Arriving after 10 minutes will be marked as **Late**.
- **80% attendance** is required.
- The course will have **Absolute Grading Scheme**.
- **Quizzes** will be **announced**. No retakes.
- **No late submissions**.
- Cheating or plagiarism will result in Zero marks.
- Taking **cross section classes** and **quizzes** are **not allowed**.
- Four Assignments and Four Quizzes. No Best off policy.

Marks distribution

- Assignments + Homework (10 %)
- Project (5 %)
- Quizzes (10 %)
- Midterms (30 %)
- Final Exam (45 %)

Books

- Stuart Russell, Peter Norvig - Artificial Intelligence A Modern Approach - (3rd Edition)
- George F. Luger - Artificial Intelligence
- Additional Reference materials:
 - Tom Mitchell - Machine Learning

Course Outline

Week	Contents
1	Introduction Agent and Agent Architectures Knowledge Representation PEAS
2	Strategies for State Space Search: Blind/Uninformed/Brute-force Search <ul style="list-style-type: none">• Depth First Search (DFS)• Breadth First Search (BFS)
3	<ul style="list-style-type: none">• Iterative Deepening Search (IDS)• Uniform Cost Search (UCS) Direction of Search, Branching Factor
4	Heuristic/Informed Search <ul style="list-style-type: none">• Best First Search• A* Search Informedness, Monotonicity, Admissibility, Optimality, Completeness
5	Adversarial Search Algorithms (Game Playing) <ul style="list-style-type: none">• Minimax Search• Alpha-Beta Pruning Evaluation Functions
Mid 1	

Quiz 1

Quiz 2

Assignment 1

Course Outline

Week	Contents	
6	Local Search: Hill Climbing Search Evolutionary Search: Genetic Algorithm <ul style="list-style-type: none">• Genetic Algorithm Introduction• Solution representation	Project Mid Evaluation
7	<ul style="list-style-type: none">• Fitness evaluation function• Selection mechanism (Roulette wheel, Rank based, Tournament)• Genetic operators (Crossover, Mutation)• Replacement method• Generation gap, Elitism	
8	Supervised learning : Regression Simple Linear Regression	Quiz 3
9	Machine Learning Introduction, Supervised Learning, Classification (ANN) <ul style="list-style-type: none">• Architecture: Feed-Forward Neural Network• Activation Functions (Linear,Sigmoid, Relu, tanh, Softamax)	
10	<ul style="list-style-type: none">• Training Algorithms: Perceptron and common activation functions, Linear separability of a perceptron• Perceptron Learning Rule• Multilayer Perceptron	Assignment 2
Mid 2		

Course Outline

Week	Contents
11	<ul style="list-style-type: none">• Back propagation algorithm (Gradient Descent Learning)• Issues of ANN (Over-fitting, Under-fitting, Validation testing)
12	<ul style="list-style-type: none">• Performance measures: Confusion Matrix (Binary & Multi Class)• Evaluation metric: Accuracy, recall, Precision, F1 score <p>Supervised learning</p> <ul style="list-style-type: none">• Decision Trees
13	Convolutional Neural Network (fwd prop only) Unsupervised learning - Clustering and its application
14	<ul style="list-style-type: none">• K-means algorithm• K-medoids algorithm
15	<ul style="list-style-type: none">• Agglomerative clustering
Finals	

Quiz 4

Assignment 3

Assignment 4

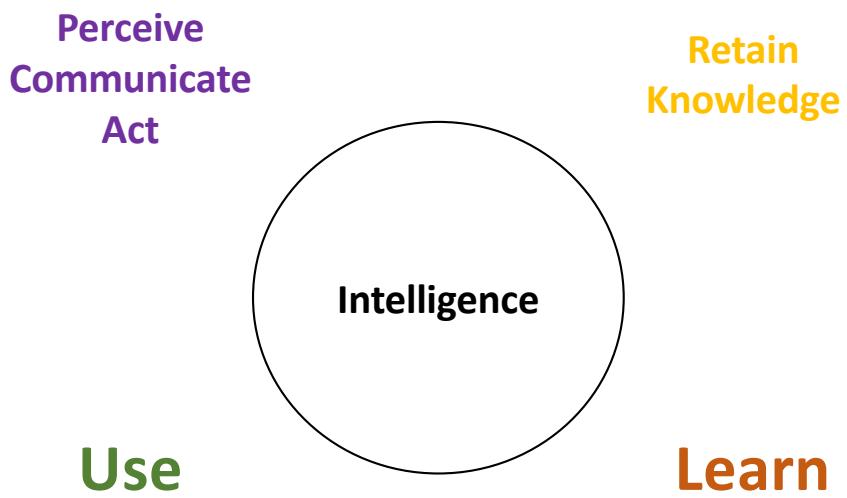
Project Final Evaluation

Introduction

Readings – Chapter 1

What is Intelligence?

- Intelligence is the ability to **acquire, understand and apply knowledge** and **skills** to **solve** problems, **adapt** to new situations, **reason** and **make** decisions.
- **Key aspects:**
 - Learning
 - Reasoning & Problem Solving
 - Adaptability
 - Memory
 - Creativity
 - Social Interactions
 - Self Awareness



Intelligence: the ability to acquire and apply knowledge and skills.

What is artificial intelligence?

Think Like a Human	Think Rationally
Act Like a Human	Act Rationally

- *The definitions on top are concerned with thought processes and reasoning,*
- *whereas the ones on the bottom address behavior.*
- *The definitions on the left measure success in terms of fidelity to human performance,*
- *Whereas the ones on the right measure against an ideal performance measure, called rationality.*

<p>Thinking Humanly</p> <p>“The exciting new effort to make computers think . . . <i>machines with minds</i>, in the full and literal sense.” (Haugeland, 1985)</p> <p>“[The automation of] activities that we associate with human thinking, activities such as decision-making, problem solving, learning . . .” (Bellman, 1978)</p>	<p>Thinking Rationally</p> <p>“The study of mental faculties through the use of computational models.” (Charniak and McDermott, 1985)</p> <p>“The study of the computations that make it possible to perceive, reason, and act.” (Winston, 1992)</p>
<p>Acting Humanly</p> <p>“The art of creating machines that perform functions that require intelligence when performed by people.” (Kurzweil, 1990)</p> <p>“The study of how to make computers do things at which, at the moment, people are better.” (Rich and Knight, 1991)</p>	<p>Acting Rationally</p> <p>“Computational Intelligence is the study of the design of intelligent agents.” (Poole <i>et al.</i>, 1998)</p> <p>“AI . . . is concerned with intelligent behavior in artifacts.” (Nilsson, 1998)</p>
<p>Figure 1.1 Some definitions of artificial intelligence, organized into four categories.</p>	

Acting humanly: The Turing Test approach

- The Turing Test, proposed by Alan Turing (1950), was designed to provide a satisfactory operational definition of intelligence.
- A computer passes the test if a human interrogator, after posing some written questions, cannot tell whether the written responses come from a person or from a computer.
- To pass the test, the computer would need to possess the following capabilities:
 - **natural language processing** to enable it to communicate successfully in English;
 - **knowledge representation** to store what it knows or hears;
 - **automated reasoning** to use the stored information to answer questions and to draw new conclusions;
 - **machine learning** to adapt to new circumstances and to detect and extrapolate patterns.
 - **computer vision** to perceive objects, and
 - **robotics** to manipulate objects and move about.

Thinking humanly: The cognitive modeling approach

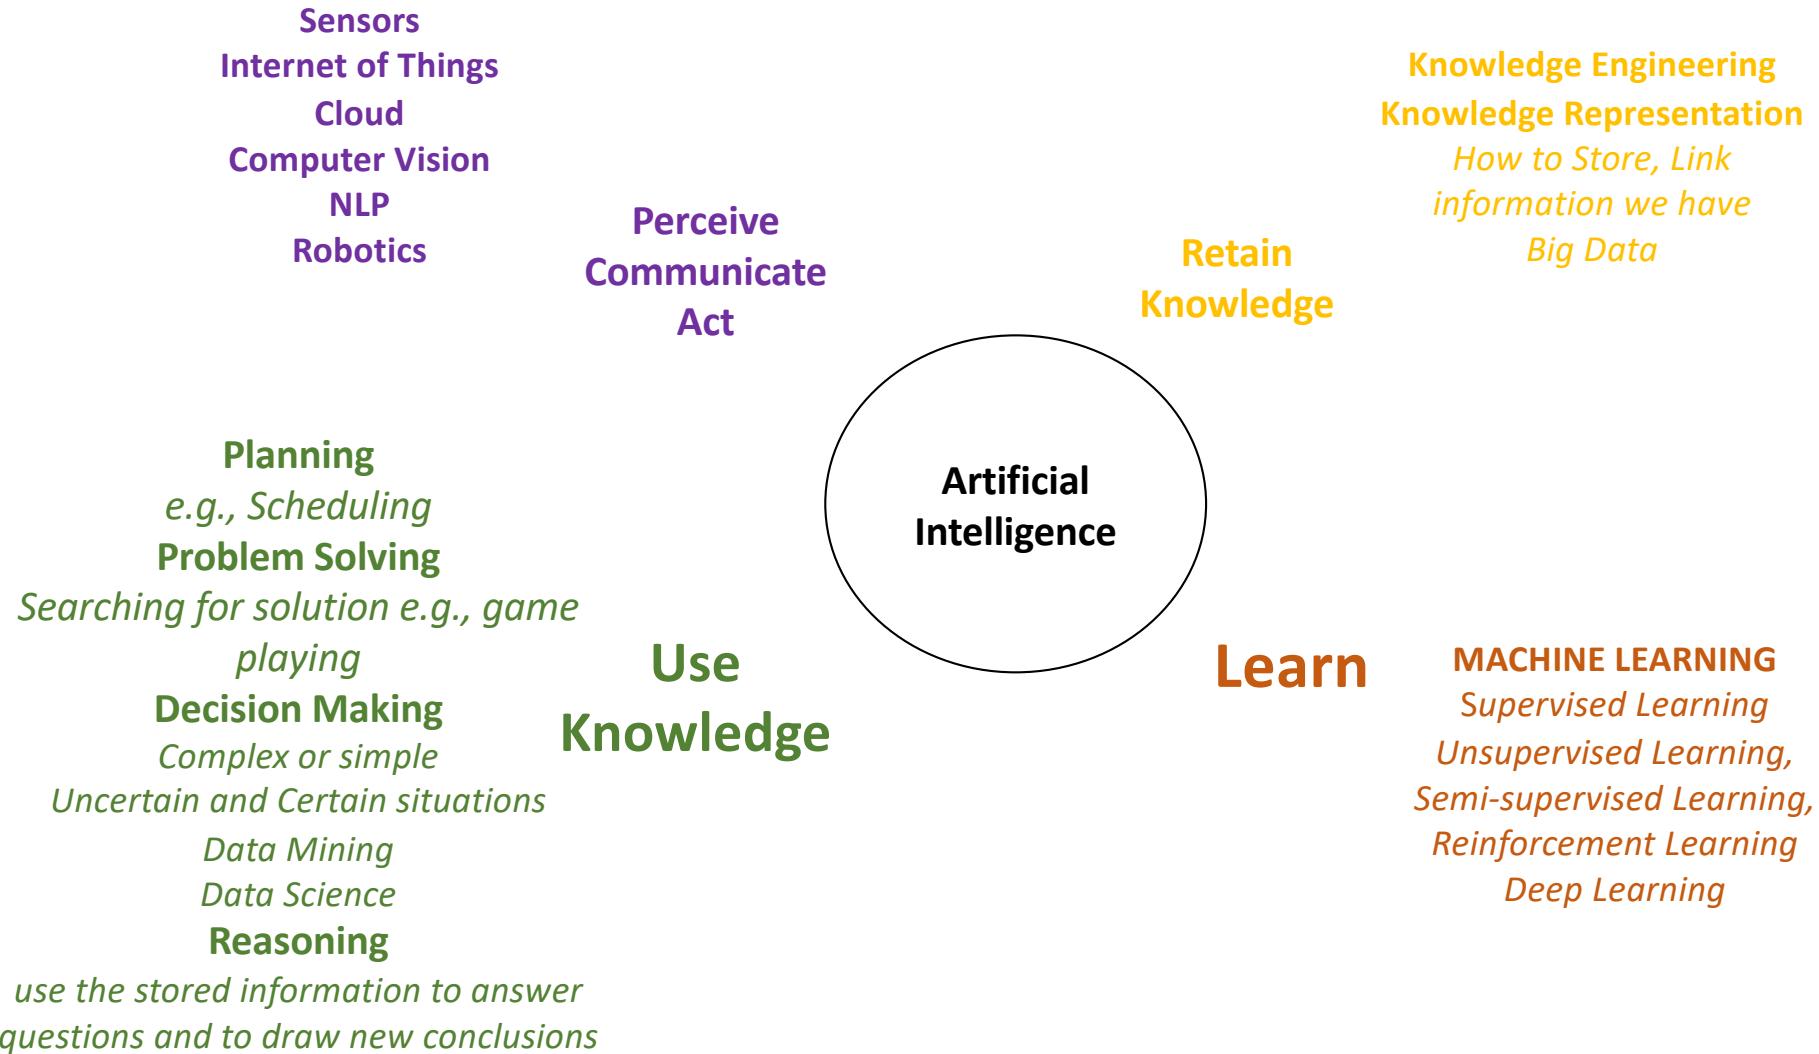
- There are three ways to do this:
 - through introspection—trying to catch our own thoughts as they go by;
 - through psychological experiments—observing a person in action;
 - and through brain imaging—observing the brain in action.
- The interdisciplinary field of cognitive science brings together computer models from AI and experimental techniques from psychology to construct precise and testable theories of the human mind.

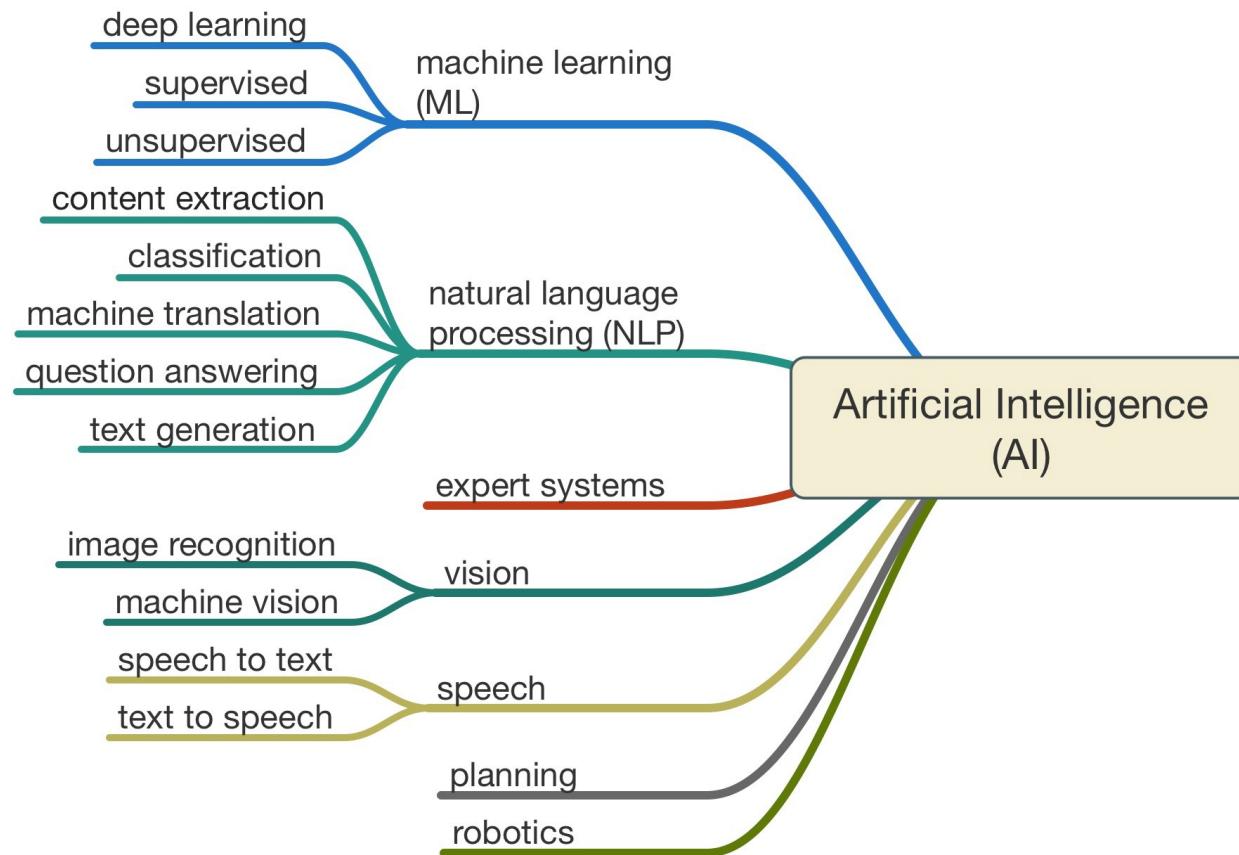
Thinking rationally: The “laws of thought” approach

- The Greek philosopher Aristotle was one of the first to attempt to codify “right thinking,” that is, irrefutable reasoning processes.
- His syllogisms provided patterns for argument structures that always yielded correct conclusions when given correct premises
- For example,
 - “Socrates is a man; all men are mortal; therefore, Socrates is mortal.”
- Programs, in principle, solve any solvable problem described in logical notation.

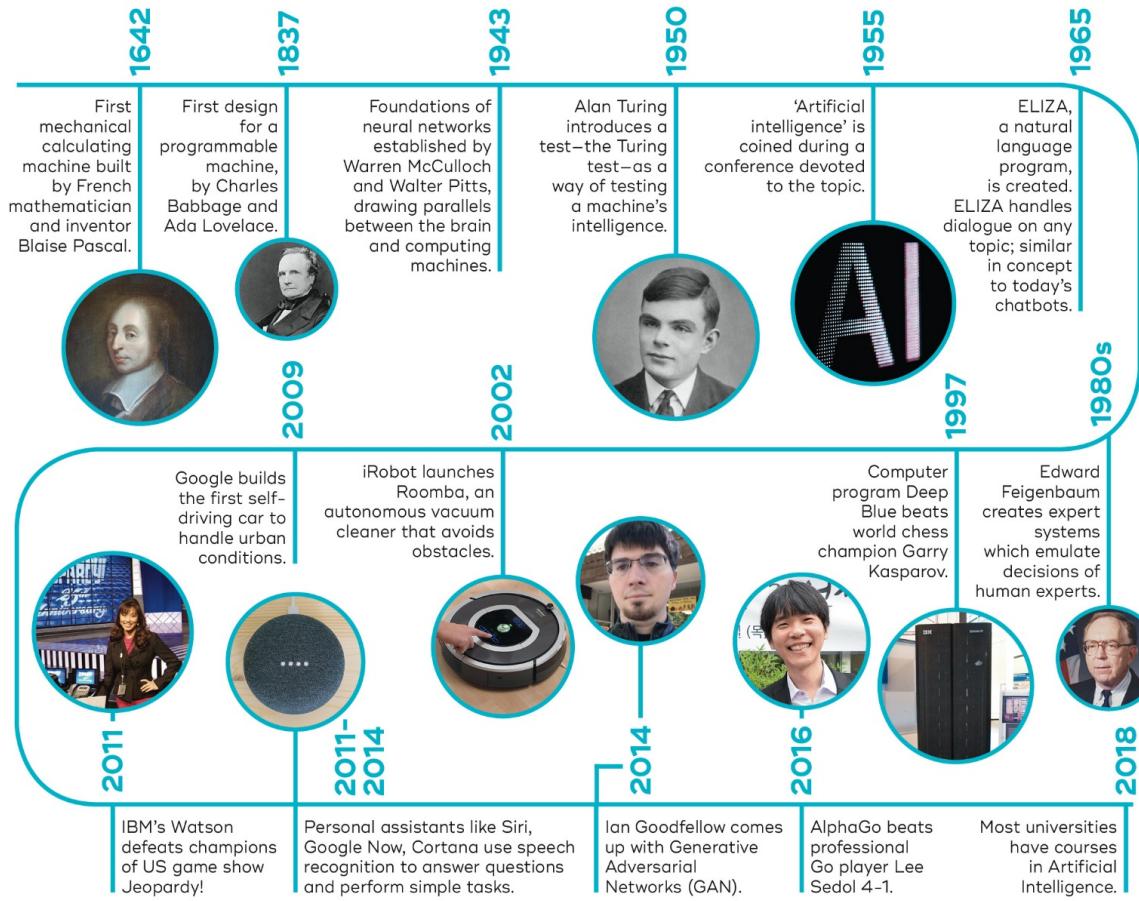
Acting rationally: The rational agent approach

- An **agent** is just something that acts.
- All computer programs do something, but computer agents are expected to do more:
 - Operate autonomously,
 - Perceive their environment,
 - Persist over a prolonged time period,
 - Adapt to change,
 - And create and pursue goals.
- A **rational agent** is one that acts so as to achieve the best outcome or, when there is uncertainty, the best expected outcome.

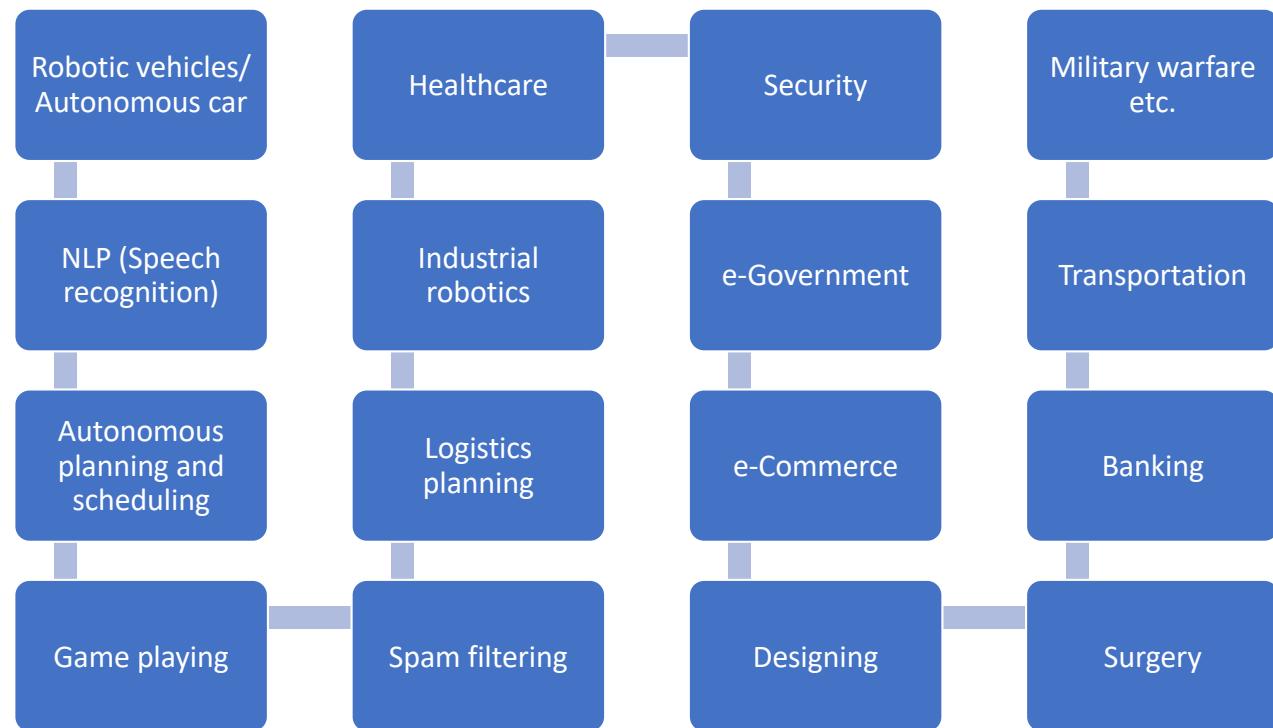




History of AI



Applications of AI



Activity

- Divide into four groups.
- Think of applications of AI and link them to a definition of AI.
- Justify your argument in five minutes.

1. AI as Thinking Like a Human

- *(Cognitive modeling approach)*
Focus: Mimicking human thought processes such as learning, reasoning, and problem-solving.
- **Applications:**
 - **Cognitive architectures** (e.g., ACT-R, SOAR) to model human thinking
 - **Intelligent tutoring systems** that adapt to a student's learning style
 - **Human memory and perception modeling** in psychology and neuroscience
 - **Brain–computer interaction systems**
 - **Simulation of human decision-making** in economics and behavioral studies

2. AI as Acting Like a Human

- *(Turing Test approach)*

Focus: Machines that behave indistinguishably from humans.

- **Applications:**

- **Chatbots and virtual assistants** (e.g., customer support bots)
- **Social robots** (e.g., humanoid robots like Sophia)
- **Speech recognition and synthesis systems**
- **Facial expression and emotion recognition**
- **Game characters and NPCs** that mimic human behavior

3. AI as Thinking Rationally

- *(Laws of thought approach)*
Focus: Logical reasoning and correct inference using formal rules.
- **Applications:**
 - Expert systems (medical diagnosis, fault detection)
 - Automated theorem proving
 - Knowledge-based systems using logic
 - Rule-based decision systems
 - Semantic reasoning and ontology-based systems

4. AI as Acting Rationally

- (*Rational agent approach – most modern AI*)
Focus: Taking actions that maximize goal achievement or performance.
- **Applications:**
 - Autonomous vehicles (self-driving cars)
 - Robotics and automation systems
 - Recommendation systems (Netflix, Amazon)
 - AI in finance (algorithmic trading, fraud detection)
 - Game-playing AI (Chess, Go, reinforcement learning agents)
 - Smart systems (smart homes, traffic control systems)

Homework

Read Chapter 1